



Introducing the Linux Kernel Programming Assignments

- ◆ All assignments implemented as kernel loadable modules.
 - ◆ NO kernel source modifications or kernel rebuilding
 - ◆ “make” with specialized kernel module Makefile
 - Access to source #include files for kernel build
 - Many kernel “helper” functions generated as in-line code from includes
 - ◆ Module inserted or removed from kernel with commands ‘insmod’ ‘rmmod’ (as root)
 - Dynamic linkage with symbols EXPORTED from kernel
 - Initialization function in module executed during insmod
 - Termination (cleanup) function in module executed during rmmod
 - ◆ Module can have both static and dynamic state variables
- ◆ Kernel loadable modules are used for device drivers, file systems, protocol stacks, etc.
 - ◆ Execute as if built directly from kernel source



Introducing the Linux Kernel Programming Assignments

- ◆ Modules loaded and executed in Red Hat Enterprise Linux 6.4 (kernel 2.6.32)
- ◆ RHEL 6.4 is to be run in virtual machine on *your* Windows, Mac OS-X, or Linux laptop/desktop
- ◆ Running in VM *is essential* because *you will* cause kernel ‘panics’ which hang the system and require a power cycle.
- ◆ The *good* news is the VM will reboot the kernel cleanly (you did not modify it permanently).
- ◆ The *bad* news is there are no effective debugging tools for the kernel so you have to be creative.



Linux 2.6.32-431 Source Space

3.4M	./tools	1.5M	./security
224K	./ipc	2.3M	./mm
4.6M	./kernel	20M	./include
40K	./usr	2.0M	./crypto
2.1M	./scripts	96K	./samples
144K	./init	1.5M	./lib
17M	./Documentation	111M	./arch (machine specific)
20M	./net	156K	./virt
19M	./sound	227M	./drivers
61M	./firmware (loaded into programmable adapters)	672K	./block
		31M	./fs
		520M	.

About 80% of source is hardware specific

COMP 530H – Fall 2014

3



Some Linux C Examples (like code you will need to understand)

```
struct pid *get_task_pid(struct task_struct *task, enum pid_type type)
{
    struct pid *pid;
    rcu_read_lock();
    if (type != PIDTYPE_PID)
        task = task->group_leader;
    pid = get_pid(task->pids[type].pid);
    rcu_read_unlock();
    return pid;
}
EXPORT_SYMBOL_GPL(get_task_pid);
```

COMP 530H – Fall 2014

4



Some Linux C Examples (like code you will need to understand)

```
pid_t pid_vnr(struct pid *pid)
{
    return pid_nr_ns(pid, task_active_pid_ns(current));
}
EXPORT_SYMBOL_GPL(pid_vnr);

pid_t pid_nr_ns(struct pid *pid, struct pid_namespace *ns)
{
    struct upid *upid;
    pid_t nr = 0;
    if (pid && ns->level <= pid->level) {
        upid = &pid->numbers[ns->level];
        if (upid->ns == ns)
            nr = upid->nr;
    }
    return nr;
}
```

COMP 530H – Fall 2014

5



Some Linux C Examples (like code you will need to understand)

```
char *d_path(const struct path *path, char *buf, int buflen)
{
    char *res;
    struct path root;
    struct path tmp;
    read_lock(&current->fs->lock);
    root = current->fs->root;
    path_get(&root);
    read_unlock(&current->fs->lock);
    spin_lock(&dcache_lock);
    tmp = root;
    res = __d_path(path, &tmp, buf, buflen);
    spin_unlock(&dcache_lock);
    path_put(&root);
    return res;
}
EXPORT_SYMBOL(d_path);
```

COMP 530H – Fall 2014

6

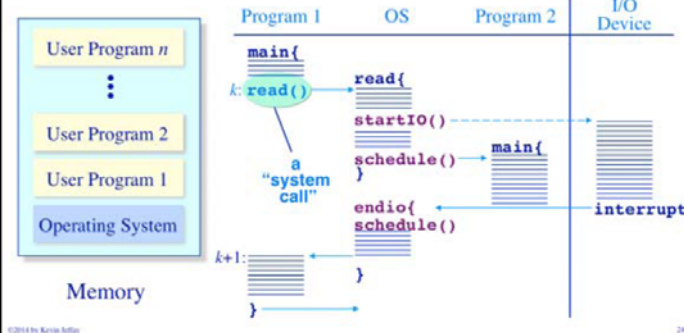


From 530 Lecture: System Calls

System Calls

The means of invoking operating system services

- Programs interact with the operating system through a procedure-call-like interface
 - Operating system procedures/functions are called *system calls*



From 530 Lectures: System Calls

System Calls & Dual-Mode Operation Mechanisms

- System calls are protected procedure calls
 - A software interrupt

```
_Read:
  LOAD r1, @SP+2
  LOAD r2, @SP+4
  LOAD r3, @SP+6
  TRAP Read_Call
```

- Software (& hardware) interrupts are used to switch from user to system mode
 - RTI ("return from interrupt") is used to switch from system to user mode

User Process

```
process UserProg
begin
:
  read(file, #bytes)
:
end UserProg

procedure Read(file, bytes)
begin
:
  Sys_Read(file, bytes)
:
end Read
```

Operating System

```
trap_handler: case trap of
  Read_Call ...
  Write_Call ...
```



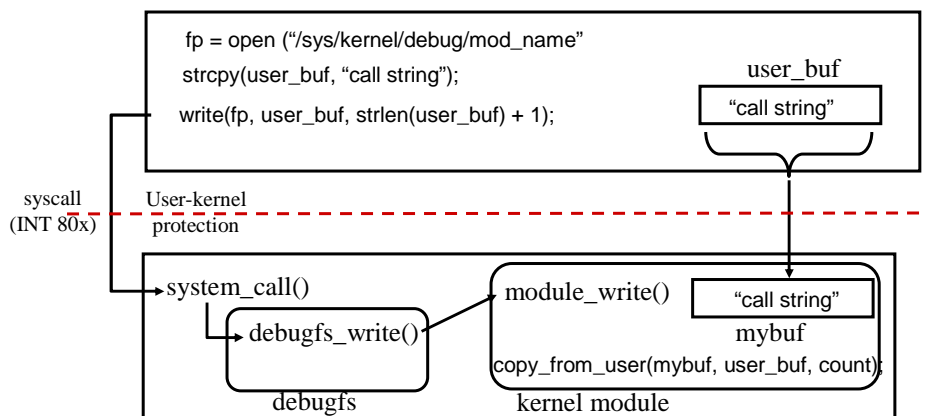
Introducing the Linux Kernel Programming Assignments

- ◆ Your modules will *emulate* a system call interface for user programs.
- ◆ A ‘pseudo’ file system, debugfs, will be used for communication between user space and kernel space.
 - ◆ User programs use read() and write() system calls
 - ◆ Kernel modules ‘hook’ the debugfs file system so reads and writes are passed to the module for execution.
- ◆ debugfs is mounted at /sys/kernel/debug
 - ◆ Kernel modules can create debugfs directories/files that can be opened and read/written by user programs
 - ◆ Kernel modules use kernel-provided functions to copy data between user space buffers and their kernel space buffers.



Introducing the Linux Kernel Programming Assignments

- ◆ Emulating a system call with debugfs (making the call)





Introducing the Linux Kernel Programming Assignments

◆ Emulating a system call with debugfs (returning result)

